

50.020 Security

Lecture 8: Operating System Security I

Introduction

Operating System Security

50.020
Security
Lecture 8:
Operating
System
Security I

We focus on Malware (i.e., Malicious software) threats to Operating System.

Focus of This Lecture:

- General things about attacks: attack stages
- Operating system attack example: buffer overflow attacks
 - How buffer overflow attacks work

Focus of Next Lecture:

- Operating system attack example: buffer overflow attacks (continue)
 - Countermeasures of buffer overflow attacks
 - Variants of buffer overflow attacks
- More about malware

Attack Stages

Attack Stages

50.020
Security
Lecture 8:
Operating
System
Security I

For any attack, the attacker needs to prepare

- Target victim identification
- Goal to achieve
- Payload that achieves the goal
- Attack vector for payload

Possible Attack Goals

A compromised machine can be used for:

- Malicious web server, distribute illegal content, phishing
- Email spam source
- Steal virtual and real money from accounts
- Launch social engineering attacks
- Bot in a botnet
- Blackmail (data encryption, publishing of private data)

Attack Strategies

50.020
Security
Lecture 8:
Operating
System
Security I

The method the attacker uses to get payload to victim

- Most common:
 - Phishing, Spearphishing
 - Social engineering
 - Watering Hole Attack
- Less common (why? NAT'ing)
 - Direct/ active attacks (on hosts)

Email Security

50.020
Security
Lecture 8:
Operating
System
Security I

- In general, emails are not secure
 - Content can be eavesdropped during transmission
 - Sender can claim any identity
 - Emails using HTML can track reads by user, contain malicious code
- Everyone knows spam, but emails can also be used for *phishing*

Phishing Attacks

- Emails that try to trick the user to visit a website
 - Pretends to be from legitimate source (e.g. Google)
 - Asks user to change the password
 - Link provided will lead to attacker website
 - Website will look like legitimate site, but steal credentials
- Often sent to random victims, easy to identify
 - User does not have account with target service
 - Typos or other unprofessional content
 - Link directs to something suspicious
- *Spearphishing* is a variant in which the attacker crafts a specific email for target, to maximise likelihood of success
 - Attacker could pretend to be a person victim knows
 - Account information of user could be included
 - Email could be sent at time where user is busy

Spearphishing Example

Your Mailbox all@sutd.edu.sg, May Shutdown in 48-72 hours.



Your Mailbox all@sutd.edu.sg, May Shutdown in 48-72 hours.

Sat 17-Feb-18 7:31 AM

Account Require Verification

Dear all,

This is **sutd.edu.sg Office Administrator** managing your account server, We're writing to inform you that your Account will be permanently locked and shutdown! on our server due to your failure to verify and re-confirm your account **ownership**.

This a secure step to keep all account updated, secured and remove all malicious treat on our server. Kindly click on www.sutd.edu.sg/re-verify to securely get verified to continue using our server.

We are going to permanently remove your account (**all@sutd.edu.sg**),
If you fail to adheld to our instruction.

NOTE: This is a one time user verification carried out in purpose to provide a more secured platform and shut down robot or malicious users created in purpose of spamming and other fraudulent activities.

Best Regards,
Administrator.

(C) 2018 sutd.edu.sg. All rights reserved. NMLSR ID 399801

Social Engineering

50.020
Security
Lecture 8:
Operating
System
Security I

- Spearphishing is an example of a *social engineering* attack
- Any attack in which the attacker tries to trick the user in performing an action
 - Spoofed emails that tell customers to send payment to different account
 - Phone calls
 - Mail
- Usually pretending to be person of authority, or in need of help
- Inherent lack of authentication in real-world is a problem

Watering Hole Attacks


50.020
Security
Lecture 8:
Operating
System
Security I

- General attack strategy: trick the user to download malware
 - Share malware in filesharing tools
 - Embed malware in websites under attacker control
- Watering Hole Attacks
 - Identify specific websites visited by victims, compromise them to inject malware

Execution of downloaded content

Malicious Attacks generated through PDF, Flash, Java, in 2013

Source: Cisco Cloud Web Security reports



- When browsing the web, lots of content is downloaded
- It is then evaluated or executed locally
 - Javascript, images, PDF, Flash, videos, ...
- Any of these is a potential attack vector!
- Modern browsers try to isolate the rendering engine from the OS

Malware

50.020
Security
Lecture 8:
Operating
System
Security I

- Lets assume attacker identified target and strategy
- A payload exists that achieves the goal if executed on victim's machine
- How to deliver the payload?
 - We need an *attack vector* to deploy payload
- Attacker vector + payload = Malware

Operating System Attack Example: Buffer Overflow Attacks

Overview Buffer Overflow

- Is a **major attack vector** to inject code to hijack the target machine's execution path. The injected code here is a form of **payload**.
- Originally exploited data/code mix in von Neumann memory
 - Data provided by attacker could be executed in place
 - Control flow data (return address) next to user data
 - Some functions allow user to write more data than intended


Overview Buffer Overflow: Continue

- Buffer overflow attacks exploit well-known vulnerabilities
 - In particular, related to C language
 - In particular, related to important insecure¹ functions in LibC
- The overflow overwrites data on the stack, which will influence control flow
 - Most importantly, the stored *return address*
 - Other variables can also be overwritten
 - Stack frame could also be manipulated

¹<https://wiki.sei.cmu.edu/confluence/display/c/MS24-C.+Do+not+use+deprecated+or+obsolescent+functions>

Refresher: Architectures

- Computers commonly use *von Neumann* architecture
 - Important for us: memory holds both *data* and *instructions*
- This fundamentally enables a series of attacks
 - Attackers are able to write data over legitimate instructions
 - Attacker's data is then executed as instructions
 - Attacker could also overwrite datastructures of running code
- Most prominent attack: *stack-based buffer overflow attacks* (our focus)
 - But also: heap overflow, ...




Von Neumann
Architecture

Memory Layout

50.020
Security
Lecture 8:
Operating
System
Security I


- Computers have different memory layers: registers, cache, RAM, discs
- Each process is presented with an abstract linear memory space to use
- OS takes care of translating memory access to the caches, RAM, disc
- Linear memory space is divided in sections
 - Stack is used for static allocation
 - Heap is used for dynamic allocation



Memory use by processes

50.020
Security
Lecture 8:
Operating
System
Security I

- Processes use stack to store data structures known at *compile-time*
 - Compiler knows size, can reserve appropriate memory area in advance
- The heap stores dynamic datastructures
 - Size or number of datastructures not known in advance
 - Memory space for that data has to be dynamically allocated (malloc)



Call Stack

- A **call stack** is a stack data structure that stores information about the active subroutines (i.e., **callees**) of a computer program (i.e., **caller**).
- A call stack can be used to keep track of the point to which each active subroutine should return control when it finishes executing.
- Watch the following videos for call stack introduction:
 - <https://www.youtube.com/watch?v=Q2sFmqvpBe0>
 - https://www.youtube.com/watch?v=XbZQ-EonR_I
(More detailed)
- Each subfunction call will add new memory space on stack for that subfunction
 - As result, stack use is growing with deeper function call nesting
- If a subfunction returns, its memory space is *freed* (not erased)
 - Next subfunction call will overwrite

Calling conventions

- Stack space for each function is called *stack frame*²
- Stack frame header is prepared by calling function
- *Calling convention* defines the way that each stack frame header is organized
- This way, caller and callee can be sure that stack is in expected format
- Here, we will only discuss the System V AMD64 ABI calling convention for Linux/Mac

²<https://www.cs.rutgers.edu/~pxk/419/notes/frames.html>

64 bit stack layout (System V AMD64)

Registers:

- RBP: base pointer, which points to the base of the stack frame
- RSP: stack pointer, which points to the top of the stack frame
- RDI, RSI, RDX, RCX, R8, R9 used to provide arguments

To call a function:

- up to 6 arguments are passed in registers
- two more arguments/pointers can be passed on the stack
- Then, the return address is stored on stack
- Then, the calling RBP is stored on stack
- RBP is set to address of old RBP
- RSP is set to RBP-(space for variables)

Example Assembly (You will know more during Lab5)

50.020
Security
Lecture 8:
Operating
System
Security I

```
[-----registers-----]
RAX: 0x7fffffffdafo --> 0x1
RBX: 0x0
RCX: 0x0
RDX: 0x7fffffffdc30 --> 0x7fffffffe076 ("LC_PAPER=en_SG.UTF-8")
RSI: 0x7fffffffdc18 --> 0x7fffffffe038 ("LC_PAPER=en_SG.UTF-8")
RDI: 0x7fffffffdafo --> 0x1
RBP: 0x7fffffffdb30 --> 0x400690 (<__libc_csu_init>: push r15)
RSP: 0x7fffffffdafo --> 0x1
RIP: 0x40064d (<main+15>: call 0x4005b6 <getlines>)
[-----code-----]
0x400642 <main+4>: sub    rsp,0x40
0x400646 <main+8>: lea    rax,[rbp-0x40]
0x40064a <main+12>: mov    rdi,rax
=> 0x40064d <main+15>: call   0x4005b6 <getlines>
```

Example minimal setup of stack for subfunction `getlines()`, and call. Call will push RBP and RIP onto stack automatically.

```
Dump of assembler code for function getlines:
0x00000000004005b6 <+0>: push    rbp
0x00000000004005b7 <+1>: mov     rbp,rsp
0x00000000004005ba <+4>: sub     rsp,0x20
...
0x00000000004005eb <+53>: leave
0x00000000004005ec <+54>: ret
End of assembler dump.
```

Subfunction `getlines()` prepares its stack frame, and `leave + ret` at end.

Example C code (8 arguments)

Example (Function call (C/AMD64/Linux))

```
long myFunc(long a, long b, long c, long d,  
            long e, long f, long g, long h)  
{  
    char myBuffer[24];  
    gets(myBuffer);  
    return a+b+c+d+e+f+g+h;  
}
```

- Important points about this function
 - $24 \cdot 8 (= 3 \cdot 64)$ bit variables
 - $8 \cdot 64$ bit arguments

Example stack layout

50.020
Security
Lecture 8:
Operating
System
Security I

AMD64 calling convention

High addr.

RBP +24

RBP +16

RBP +8

RBP

RBP -8

RBP -16

RSP

Low addr.

calling frame

h

g

stored RIP

saved RBP

myBuffer[16]

myBuffer[8]

myBuffer[0]

Registers

RDI

a

RSI

b

RDX

c

RCX

d

R8

e

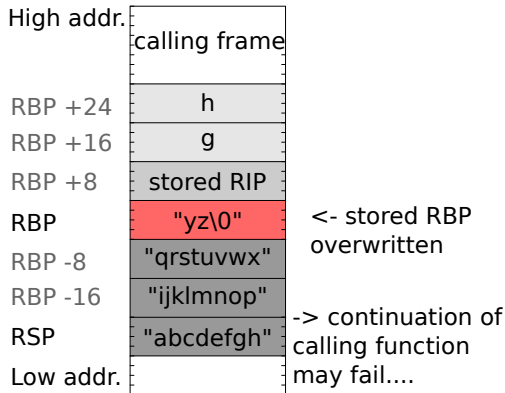
R9

f

Note: Stack
"grows downwards"

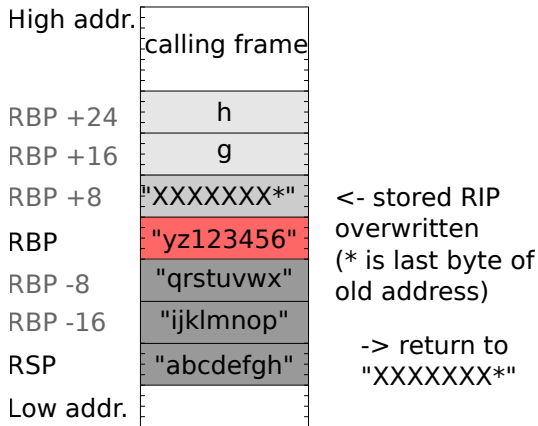
Simple Buffer Overflow

- Assume myBuffer contains
abcdefghijklmnopqrstuvwxy\0
- gets() continues to read in until \0 char



Malicious Buffer Overflow

- Assume myBuffer contains
 abcdefghijklmnopqrstuvwxyz123456XXXXXXX\0
- return addr. is overwritten with XXXXXXXX



Code injection attacks

50.020
Security
Lecture 8:
Operating
System
Security I

- The return address would point to the address of some malicious code designed by Attacker
- If stack is executable:
 - Easy way to execute arbitrary code ("shellcode")
 - For example, shellcode can spawn reverse shell to attacker

Next Lecture

50.020
Security
Lecture 8:
Operating
System
Security I

- Continue with buffer overflow attacks:
 - Countermeasures of buffer overflow attacks
 - Variants of buffer overflow attacks
- More about Malware